

Linear Regression


```
In [1]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures # "Feature Engineering" for
from sklearn.metrics import mean_squared_error
```

```
In [2]: %matplotlib inline
%config InlineBackend.figure_formats = ['svg']
```

Assume that the (unknown) target function is

$$y = f(x) = 1.0 + 0.5 \times x$$

```
In [3]: # Create some training data item
np.random.seed(42)
N = 10
X_train = 3.0*np.random.rand(N, 1) # rand(N,1): Create N random numbers in the i
y_train = 1.0 + 0.5*X_train # randn(N,1): Create N random numbers ~ standard no
plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', markersize=10, label='train')
plt.legend()
plt.show()
```




The data items generated above don't have any noise, so they all lie on a straight line.

To simulate the fact that data collected in practice usually contain certain noises, we add some noises to the generated target values.

$$y = f(x) = 1.0 + 0.5 \times x + \varepsilon$$

```
In [4]: # Create some training data item
np.random.seed(42)
N = 10
X_train = 3.0*np.random.rand(N, 1) # rand(N,1): Create N random numbers in the i
y_train = 1.0 + 0.5*X_train + np.random.randn(N,1)/2.0 # randn(N,1): Create N r
```

```
In [5]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', markersize=10, label='train')
#plt.plot(X_test, y_test, 'r*', markersize=10, label='test')
plt.legend()
plt.show()
```




```
In [6]: # Create some data to use for plotting
X_new = np.linspace(0,3,100).reshape(100,1) # Create an array of 100 equal-space
```

```
In [7]: lin_reg = LinearRegression() # Normal Equation
lin_reg.fit(X_train,y_train)

plt.figure(figsize=(6,6))
plt.plot(X_train,y_train,'b.', ms=10, label='train')
plt.plot(X_new, lin_reg.predict(X_new), 'k-')
```

```
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()
print(lin_reg.intercept_, lin_reg.coef_) # intercept_:  $w_0$  (sometimes we call it  $b$ )
# Here we only have a simple linear regression model:  $\text{predict} = w_0 + w_1 \cdot x$ 

y_pred_train = lin_reg.predict(X_train) # Use the trained model to make predictions
lr_loss_train = mean_squared_error(y_train, y_pred_train) # Mean Squared Error Loss
print('Training Loss: ', lr_loss_train)
```



```
[0.33181066] [[0.75155622]]
Training Loss: 0.09991531905633781
```

```
In [8]: print(X_train)
        print(y_train)
```

```

[[1.12362036]
 [2.85214292]
 [2.19598183]
 [1.79597545]
 [0.46805592]
 [0.46798356]
 [0.17425084]
 [2.59852844]
 [1.80334504]
 [2.12421773]]
[[1.32707299]
 [2.69735148]
 [1.86628207]
 [1.66512285]
 [1.3550091 ]
 [0.27735166]
 [0.2246665 ]
 [2.01812045]
 [1.39525696]
 [2.21923253]]

```

After training, we would like to evaluate the trained model on "test" data.

```

In [9]: # Create some test data
N_test = 10
X_test = 3.0*np.random.rand(N_test, 1)
y_test = 1.0 + 0.5*X_test + np.random.randn(N_test,1)/2.0 # Note: Testing dataset

```


```

In [10]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10, label='train')
plt.plot(X_new, lin_reg.predict(X_new), 'k-')
plt.plot(X_test, y_test, 'r*', ms=10, label='test')
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()

y_pred_train = lin_reg.predict(X_train) # Use the trained model to make prediction
lr_loss_train = mean_squared_error(y_train, y_pred_train) # Mean Squared Error Loss
print('Training Loss: ', lr_loss_train)

y_pred_test = lin_reg.predict(X_test) # Use the trained model to make prediction
lr_loss_test = mean_squared_error(y_test, y_pred_test) # Mean Squared Error Loss
print('Testing Loss: ', lr_loss_test)

```



Training Loss: 0.09991531905633781

Testing Loss: 0.21020128681013045

Polynomial Regression

$$\text{prediction} = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_Dx^D$$


```
In [11]: deg = 3
poly = PolynomialFeatures(degree=deg, include_bias=False)
X_poly = poly.fit_transform(X_train) # x --> x, x^2, x^3
X_new_poly = poly.fit_transform(X_new) # x --> x, x^2, x^3

poly_reg = LinearRegression() # In essence, Polynomial Regression is still Linear Regression
poly_reg.fit(X_poly, y_train)

plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10, label='data')
# plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
# print(lin_reg.intercept_, lin_reg.coef_)

plt.plot(X_new, poly_reg.predict(X_new_poly), 'g', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()
poly_loss = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', poly_loss)
```

```
[-0.0916451] [[ 2.35120734 -1.23811156  0.26177074]]
```



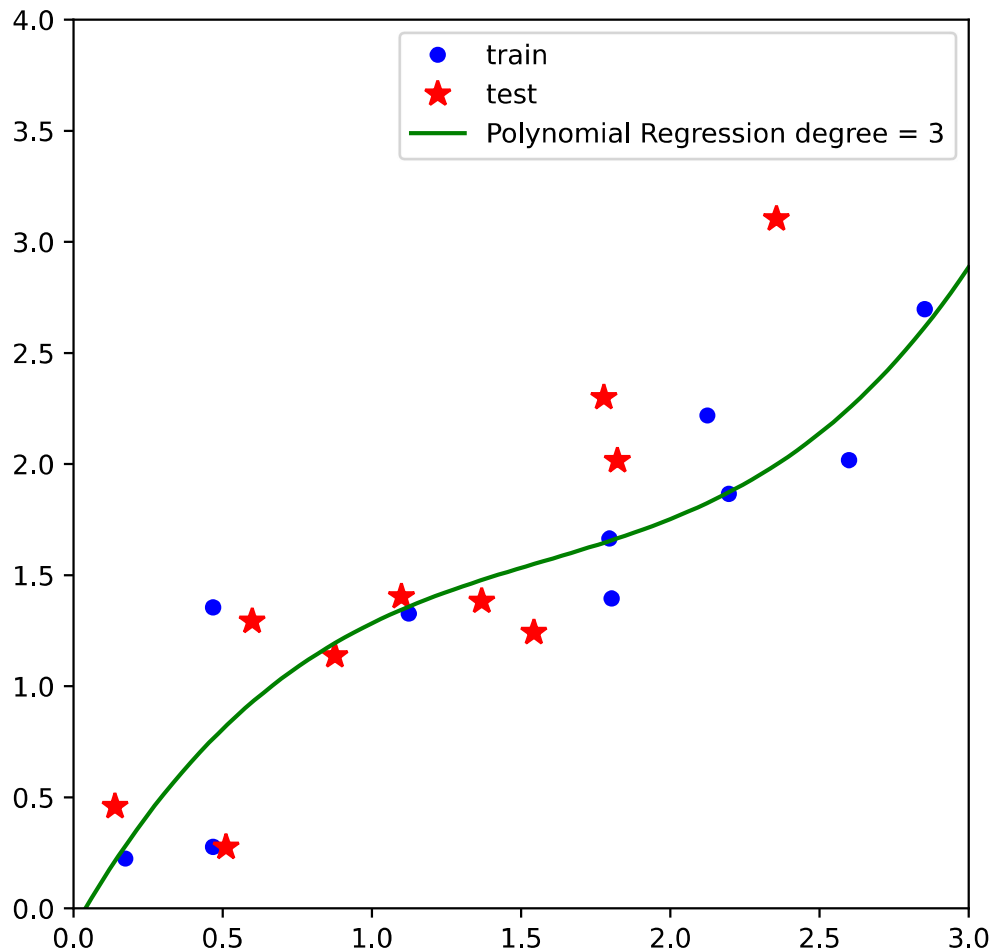
Training Loss: 0.0875019006667126

```
In [12]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10, label='train')
plt.plot(X_test, y_test, 'r*', ms=10, label='test')
# plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
# print(lin_reg.intercept_, lin_reg.coef_)

plt.plot(X_new, poly_reg.predict(X_new_poly), 'g', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()

poly_loss_train = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', poly_loss_train)
X_test_poly = poly.fit_transform(X_test)
poly_loss_test = mean_squared_error(y_test, poly_reg.predict(X_test_poly))
print('Testing Loss: ', poly_loss_test)
```

```
[-0.0916451] [[ 2.35120734 -1.23811156  0.26177074]]
```



Training Loss: 0.0875019006667126

Testing Loss: 0.23747048298743464

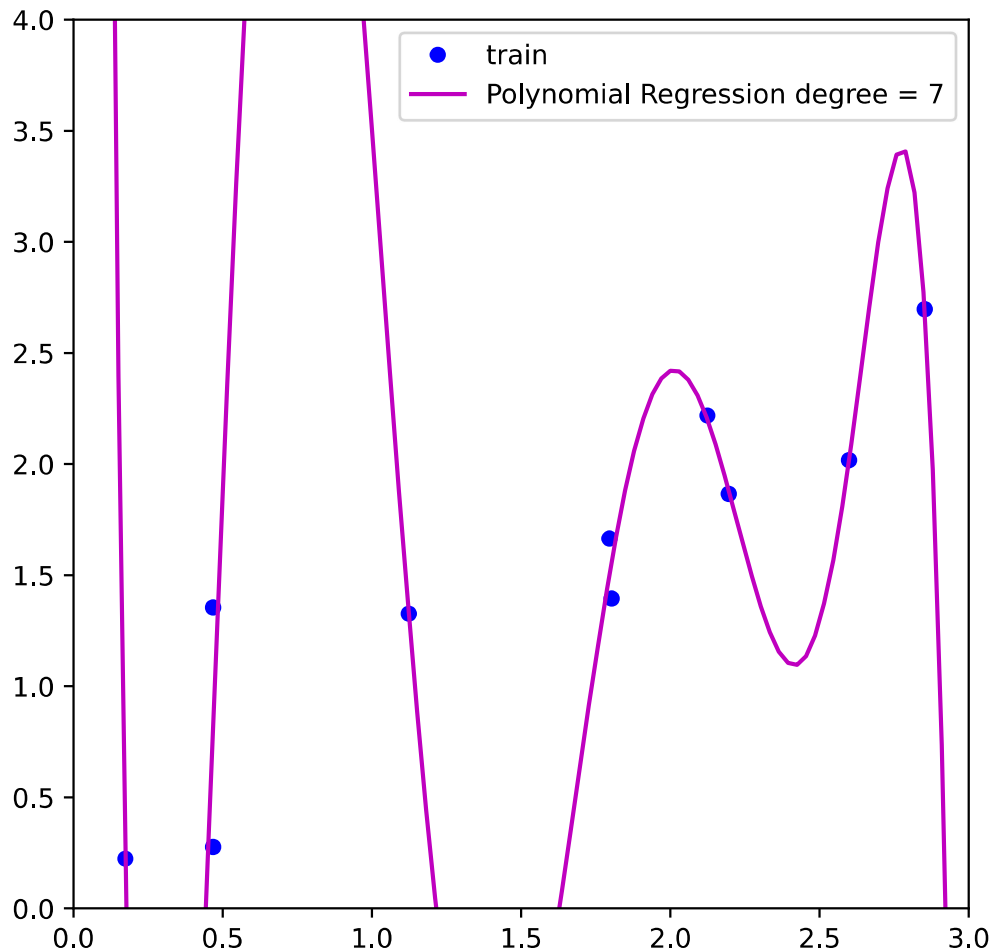
```
In [13]: deg = 7
poly = PolynomialFeatures(degree=deg, include_bias=False)
X_poly = poly.fit_transform(X_train) # x --> x, x^2, x^3, ..., x^7
X_new_poly = poly.fit_transform(X_new) # x --> x, x^2, x^3, ..., x^7

poly_reg = LinearRegression() # In essence, Polynomial Regression is still Linear
poly_reg.fit(X_poly, y_train)

plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10, label='train')
# plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
# print(lin_reg.intercept_, lin_reg.coef_)

plt.plot(X_new, poly_reg.predict(X_new_poly), 'm', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()
poly_loss_train = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', poly_loss_train)
```

```
[38.069871] [[ -399.60873887  1364.69579468 -2093.21939506  1665.82047777
 -716.24860915   158.03695709  -14.0390417  ]]
```



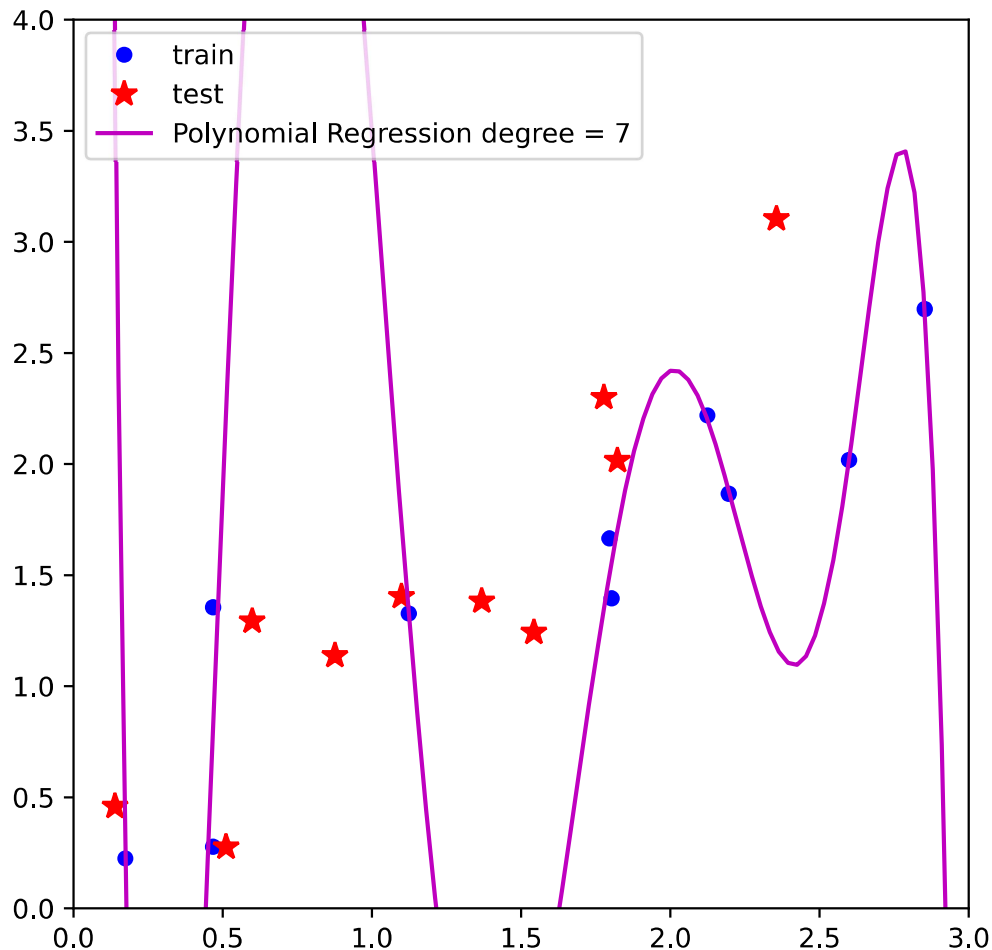
Training Loss: 0.06322418953679351

```
In [14]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10, label='train')
plt.plot(X_test, y_test, 'r*', ms=10, label='test')
# plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
# print(lin_reg.intercept_, lin_reg.coef_)

plt.plot(X_new, poly_reg.predict(X_new_poly), 'm', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
plt.legend()
plt.axis([0, 3, 0, 4])
plt.show()

poly_loss_train = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', poly_loss_train)
X_test_poly = poly.fit_transform(X_test)
poly_loss_test = mean_squared_error(y_test, poly_reg.predict(X_test_poly))
print('Testing Loss: ', poly_loss_test)
```

```
[38.069871] [[ -399.60873887  1364.69579468 -2093.21939506  1665.82047777
 -716.24860915   158.03695709  -14.0390417  ]]
```

Training Loss: 0.06322418953679351

Testing Loss: 5.957556908635877

```
In [15]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10)
plt.plot(X_test, y_test, 'r*', ms=10)
plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
print('---Simple Linear Regression---')
print(lin_reg.intercept_, lin_reg.coef_)
lr_loss_train = mean_squared_error(y_train, lin_reg.predict(X_train))
print('Training Loss: ', lr_loss_train)
lr_loss = mean_squared_error(y_test, lin_reg.predict(X_test))
print('Testing Loss: ', lr_loss)
print()

deg = 3
print(f'---Polynomial Regression degree = {deg}---')
poly = PolynomialFeatures(degree=deg, include_bias=False)
X_poly = poly.fit_transform(X_train) # x --> x, x^2, x^3
X_new_poly = poly.fit_transform(X_new) # x --> x, x^2, x^3

poly_reg = LinearRegression() # In essence, Polynomial Regression is still Line
poly_reg.fit(X_poly, y_train)
plt.plot(X_new, poly_reg.predict(X_new_poly), 'g', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
lr_loss_train = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', lr_loss_train)
X_test_poly = poly.fit_transform(X_test)
poly_3_loss = mean_squared_error(y_test, poly_reg.predict(X_test_poly))
print('Testing Loss: ', poly_3_loss)
print()
```

```

deg = 7
print(f'---Polynomial Regression degree = {deg}---')
poly = PolynomialFeatures(degree=deg, include_bias=False)
X_poly = poly.fit_transform(X_train) # x --> x, x^2, x^3, ..., x^7
X_new_poly = poly.fit_transform(X_new) # x --> x, x^2, x^3, ... x^7

poly_reg = LinearRegression()
poly_reg.fit(X_poly, y_train)
plt.plot(X_new, poly_reg.predict(X_new_poly), 'm', label='Polynomial Regression')
print(poly_reg.intercept_, poly_reg.coef_)
lr_loss_train = mean_squared_error(y_train, poly_reg.predict(X_poly))
print('Training Loss: ', lr_loss_train)
X_test_poly = poly.fit_transform(X_test)
poly_7_loss = mean_squared_error(y_test, poly_reg.predict(X_test_poly))
print('Testing Loss: ', poly_7_loss)

plt.axis([0, 3, 0, 4])
plt.legend()
plt.show()

```

---Simple Linear Regression---

[0.33181066] [[0.75155622]]

Training Loss: 0.09991531905633781

Testing Loss: 0.21020128681013045

---Polynomial Regression degree = 3---

[-0.0916451] [[2.35120734 -1.23811156 0.26177074]]

Training Loss: 0.0875019006667126

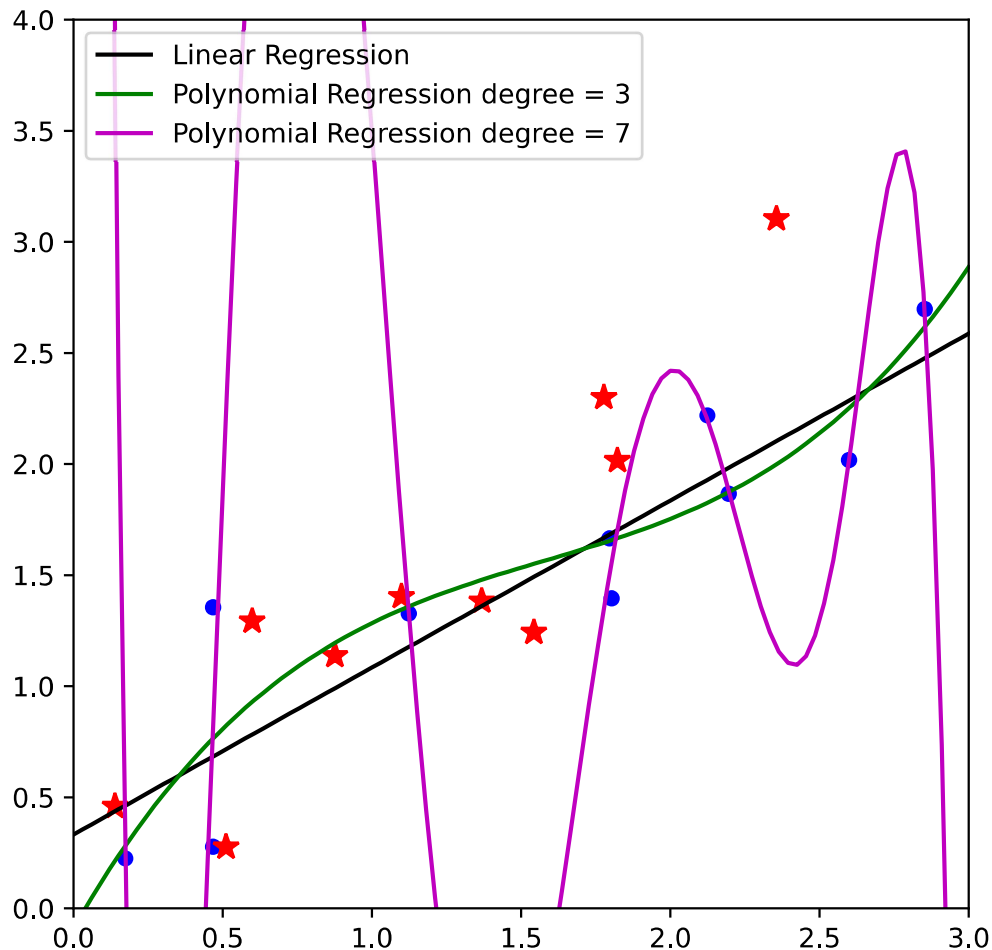
Testing Loss: 0.23747048298743464

---Polynomial Regression degree = 7---

[38.069871] [[-399.60873887 1364.69579468 -2093.21939506 1665.82047777
-716.24860915 158.03695709 -14.0390417]]

Training Loss: 0.06322418953679351

Testing Loss: 5.957556908635877



The 7-th degree polynomial regression model above (magenta color) is overfit to our training dataset. If we use an overfit model on our future data, it will have a lot of errors.

The magnitude of parameters of an overfit model would be very big (very far from 0.0).

Ridge Regression

Ridge Regression: Weight Decay/Maximum A Posteriori Estimation (MAP) for Linear Regression

prediction = $w_0 + w_1x + w_2x^2 + \dots + w_Dx^D$ (However w_1, w_2, \dots, w_D will have small values close to 0.0).

$$w_{map} = \arg \min_w NLL(w) + \lambda \|w\|_2^2$$

$\lambda \|w\|_2^2$: Regularization term, and λ is the regularization strength.

```
In [16]: plt.figure(figsize=(6,6))
plt.plot(X_train, y_train, 'b.', ms=10)
plt.axis([0, 3, 0, 4])

deg = 7
poly = PolynomialFeatures(degree=deg, include_bias=False)
X_poly = poly.fit_transform(X_train) # x --> x, x^2, x^3, ..., x^7
X_new_poly = poly.fit_transform(X_new)
```

```

plt.plot(X_new, lin_reg.predict(X_new), 'k-', label='Linear Regression')
print(lin_reg.intercept_, lin_reg.coef_)

lambda = 0
ridge_reg = Ridge(alpha=lambda)
ridge_reg.fit(X_poly, y_train)
plt.plot(X_new, ridge_reg.predict(X_new_poly), 'm-', label=r'$\lambda=\{\}$'.format(lambda))
print(ridge_reg.intercept_, ridge_reg.coef_)

lambda = 10**-5 # 10^{-5}
ridge_reg = Ridge(alpha=lambda)
ridge_reg.fit(X_poly, y_train)
plt.plot(X_new, ridge_reg.predict(X_new_poly), 'g-', label=r'$\lambda=\{\}$'.format(lambda))
print(ridge_reg.intercept_, ridge_reg.coef_)

lambda = 1000
ridge_reg = Ridge(alpha=lambda)
ridge_reg.fit(X_poly, y_train)
plt.plot(X_new, ridge_reg.predict(X_new_poly), 'b-', label=r'$\lambda=\{\}$'.format(lambda))
print(ridge_reg.intercept_, ridge_reg.coef_)

lambda = 10000000000 # Choose it by our experience, trial and error
ridge_reg = Ridge(alpha=lambda)
ridge_reg.fit(X_poly, y_train)
plt.plot(X_new, ridge_reg.predict(X_new_poly), 'r-', label=r'$\lambda=\{\}$'.format(lambda))
print(ridge_reg.intercept_, ridge_reg.coef_)

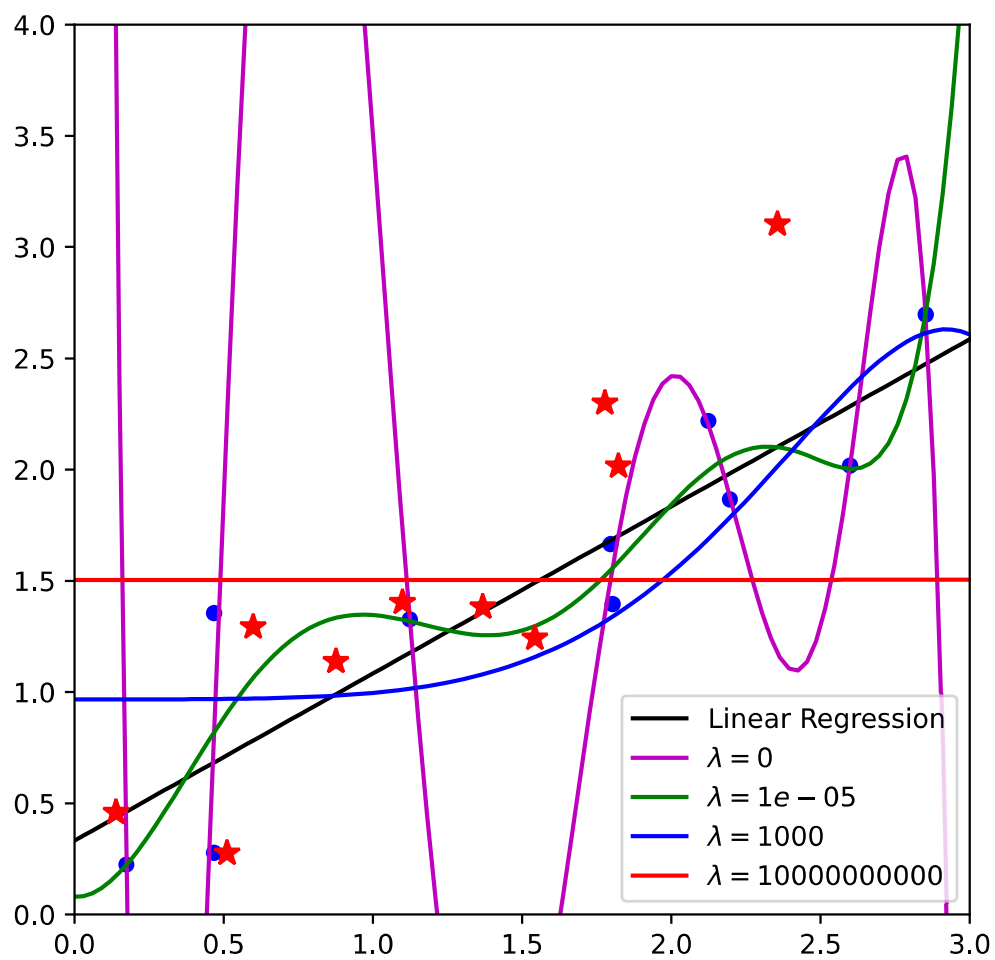
plt.plot(X_test, y_test, 'r*', ms=10)
plt.legend()
plt.show()

```

```

[0.33181066] [[0.75155622]]
[38.06834626] [[ -399.59270239  1364.64125612 -2093.13554384  1665.75343429
 -716.21961323   158.03051773  -14.03846577]]
[0.07987999] [[-0.09370931  6.14945012 -4.77478402 -3.78021829  5.75140457 -2.283
94658
 0.29814782]]
[0.96643294] [[ 0.00141423  0.00266913  0.00442294  0.00683627  0.00931761  0.008
96911
 -0.00359498]]
[1.50446301] [[6.07278017e-10 1.77351966e-09 4.72471362e-09 1.25448975e-08
3.35844395e-08 9.07463757e-08 2.47260904e-07]]

```



In [16]: